

Submission for IPDPS 2018 PhD forum

*ECE, Rutgers University
§ Brookhaven National Laboratory

I. APPLICANT STATEMENT OF INTEREST

A. Short bio sketch

I am currently a third year PhD candidate in the Electrical and Computer Engineering (ECE) department at Rutgers University and a member of the RADICAL group, lead by Dr. Shantenu Jha. Prior to my PhD program, I obtained my Master's in ECE working with the RADICAL group at Rutgers University and a Bachelor's in Electronics and Communication Engineering from SSN College of Engineering, Anna University, Chennai, India.

B. Current research interests

I was introduced to the concept of abstractions and ensemble-based applications in molecular dynamics during my research, towards my Master's, with the RADICAL group.

C. Plan for future research/career

D. Objectives for participating in IPDPS Student Program

- How do you expect the interaction with the IPDPS community will further your research/career goals?

Comments for reviewers were very helpful.

II. POSTER PROPOSAL

A. Title of work and author and advisor(s) names and affiliations

The title of the poster is “Harnessing the Power of Many: Extensible Toolkit for Scalable Ensemble Applications”. The author presenting the poster will be Vivek Balasubramanian, Electrical and Computer Engineering, Rutgers, the State University of New Jersey, under the guidance of Dr. Shantenu Jha and Dr. Matteo Turilli, Electrical and Computer Engineering, Rutgers, the State University of New Jersey.

Many scientific problems solved on HPCs consist of applications that rely on the collective output of multiple tasks as opposed to a single large task. The individual tasks within the collection might be uncoupled; if coupled, the tasks might have global or local synchronizations, and regular or irregular communication. This is in contrast to traditional parameter sweeps, or high-throughput computing (HTC) applications, where the tasks are typically identical, uncoupled, idempotent and can be executed in any order.

The execution of an ensemble on HPC machines presents three main challenges: (1) encoding scientific problems into algorithms that are amenable to distributed and coordinated solution; (2) sizing, acquiring, and managing resources for the execution; and (3) managing the execution of the ensemble.




Fig. 1: EnTK prototype

In the absence of generic solutions to these challenges, we designed and implemented the Ensemble Toolkit. EnTK adheres to the building blocks approach [] and thus supports coupling with existing runtime systems. In addition to extensibility, EnTK enables composability by providing components specifically to create ensemble-based applications. These features enable EnTK to overcome flexibility constraints of monolithic workflow systems, and thereby support a wide range of application requirements.

EnTK is engineered for scale and diversity of computing platforms and runtime systems. EnTK addresses these challenges by decoupling the description of ensemble-based applications from their execution into three separate orders of concern: specification of task and resource requirements; resource selection and acquisition; and task execution management.

B. Abstract of results to be presented in poster

We present three sets of results in the poster. We first present the performance benchmark of an EnTK prototype that provide a reference hardware configuration to support execution of up to $O(10^6)$ tasks. In the second set of results, we present the weak and strong scaling behavior of EnTK. Lastly, we implement and execute at scale our two use cases, seismic inversion and adaptive analog workflow.

We prototyped EnTK instantiating only multiple producers and consumers of tasks. The producers push into RabbitMQ queues and consumers pull from these queues. With a total of 10^6 tasks, we benchmark the prototype to observe the total execution time, base and peak memory consumption as a function of the number of producers, consumers and queues.

In Figure 1, we show that the execution duration decreases linearly at the cost of increased memory usage. We also gather from the benchmark that EnTK can be tuned, by varying the number of producers, consumers and queues, depending on the resource and application requirements.




Fig. 2: Weak scaling




Fig. 3: Strong scaling

Next we present experiments to characterize weak and strong scalability of EnTK. Weak scaling experiments describe the performance of EnTK when supporting varying levels of concurrency; Strong scaling experiments describe the performance of EnTK when executing a workload larger than the amount of available resources. For the purposes of the poster, we focus only on the Task Execution Time and the EnTK Management Overhead. The remaining durations are a function of the Python language or the underlying RTS. Enhancements to these will directly reflect improvements in these durations and thus we consider them out of scope.

In the case of weak scaling, we run four applications each with 1 pipeline, 1 stage per pipeline, 512, 1024, 2048 and 4096 tasks per stage and number of cores proportional to the number of tasks in the stage. While in the case of strong scaling, we run three applications each with 1 pipeline, 1 stage per pipeline, 8192 tasks per stage and 1024, 2048 and 4096 cores. Each task, configured to run on 1 core for ≈ 600 seconds, executes a Gromacs simulation. Each task requires an input of 4 files with an aggregate size of 550KB.

In Figure 2, the Task Execution Time increases even though there are sufficient resources to run all the tasks. We attribute this to the limitations in the current implementation of the RTS and the ORTE distributed virtual machine of OpenMP. EnTK Management Overhead remains constant up to 2048 above which the number of tasks strain the resources leading to an increase in the overhead. The Data Staging Duration increases since the total amount of data increases with increase in the number of tasks.

In Figure 3, the Task Execution Time reduces linearly with an increase in the number of available cores. EnTK Management Overhead remains constant at ≈ 8 seconds and Data Staging Duration remains constant at ≈ 90 seconds since the total number of tasks remains constant and thus the total data




Fig. 4: Specfem execution time and failure rate

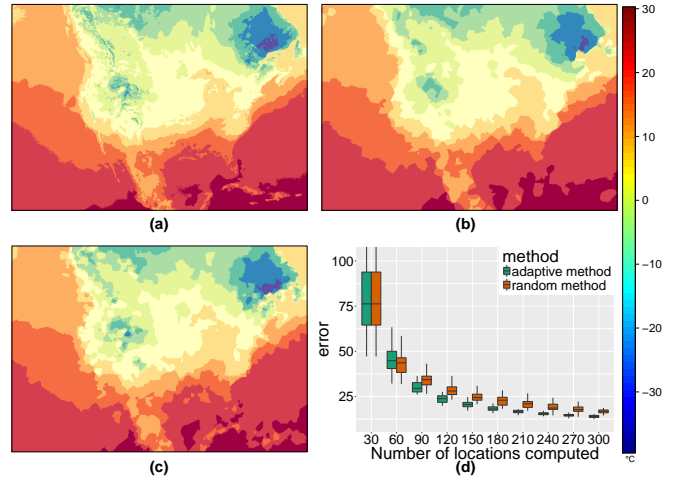


Fig. 5: Error in Analog Ensemble execution

staged is also constant.

We use EnTK to encode the forward simulations of the seismic tomography workflow. These simulations account for more than 90% of the computation time of the workflow, requiring 384 nodes for each earthquake simulation, and 40MB of input data each. Concurrent simulation of earthquakes incur a high failure rate and are resubmitted automatically by EnTK.

In Figure 4, we run six applications with 1 pipeline, 1 stage per pipeline and 1, 2, 4, 8, 16 and 32 tasks per stage. Each of the applications is run at different levels of concurrency. We observe that for constant concurrency, the execution time increases linearly with increase in the number of tasks and the execution time reduces linearly when increasing the concurrency. We also observe that reducing concurrency eliminates failures: we encountered no failures in executions with up to $O(2^4)$ concurrent tasks and 6,144 nodes. At $O(2^5)$ concurrent tasks and 12,288 nodes, 157 tasks failed. EnTK automatically submits a failed tasks and as a result we observed a Task Execution Time of ≈ 360 seconds.

We use EnTK to implement the AUA algorithm to iteratively and dynamically identify locations of the analogs. We perform experiments to compare our implementation with the status quo method, i.e. random selection of locations, and observe the speedup of the proposed algorithm.

Figure 5 shows the prediction maps and errors obtained from the two implementations. The AUA algorithm generates a map (b, c) with certain areas that have a better representation of

the analysis than the map generated by a random selection of pixels. Box plot (d) shows that the error converges faster in the AUA algorithm than in the random selection.